

Presidential Forecasting with bigKRLS

Pete Mohanty, Stanford University

Robert B. Shaffer, University of Texas at Austin

BARUG October 2016

Overview

- ▶ **Part I:** Kernel Regularized Least Squares (KRLS)
 - ▶ What is KRLS? Why use it?
 - ▶ Speed/memory tradeoffs and bigKRLS
- ▶ **Part II:** Political Forecasting and bigKRLS
 - ▶ Exploratory Modeling in the Two-Party Case
 - ▶ (The Limits of) Pooling Publicly Available Polls
 - ▶ Forecasting Voting Choices

Kernel Regularized Least Squares

Motivations

- ▶ Maximize inference and out-of-sample prediction, minimize assumptions
- ▶ estimate not just “average” but “actual” marginal effects

Statistical Properties & Applications

- ▶ For details on statistical properties, see [Hainmueller & Hazlett 2013](#)
- ▶ “Kernel balancing” may help observational studies approximate experiments; see, e.g., [Hazlett](#); [Hastie et al 2008](#)

The Challenge

Speed

- ▶ Luke Sonnet notes speed problems with *KRLS* starting at $N \approx 1,000$ but has made substantial speed gains in [Julia](#).

Size

- ▶ “Tikhonov regularization requires computation of weight matrices of dimension $\mathbf{N} \times \mathbf{N}$ which [...] may be unsuitable for large datasets.” ([Racine and Hayfield](#))
- ▶ Typical machines hit “cannot allocate vector” limits at $N \approx 3,000$ with *KRLS*

Some details

We assume that the objective function $\mathbf{y} = \mathbf{f}(\mathbf{x})$ can be approximated by

$$\mathbf{y} = \mathbf{K}\mathbf{c}$$

With \mathbf{K} the Gaussian kernel,

$$\mathbf{K} = e^{-||\mathbf{x}_i - \mathbf{x}_j||^2 / \sigma^2}$$

and \mathbf{c} a vector of weights chosen based on an L_2 penalty,

$$\hat{\mathbf{c}} = \underset{\mathbf{c} \in \mathbb{R}^P}{\operatorname{argmin}} (\mathbf{y} - \mathbf{K}\mathbf{c})'(\mathbf{y} - \mathbf{K}\mathbf{c}) + \lambda \mathbf{c}'\mathbf{K}\mathbf{c}$$

Introducing bigKRLS

bigKRLS: a new version of the algorithm which minimizes memory constraints and boosts speed in **R**.

- ▶ Reduce peak memory requirements from $\approx 9PN^2$ to $\approx 5N^2$
- ▶ “big R” ([bigmemory](#), [bigalgebra](#) & [biganalytics](#))
- ▶ $R \rightarrow C++$ ([Rcpp](#) & [RcppArmadillo](#))

bigKRLS Complexity

	Major Steps	Runtime	Memory
(1)	Standardize $\mathbf{X}_{N \times P}$, \mathbf{y}	—	—
(2)	Calculate kernel $\mathbf{K}_{N \times N}$	$O(N^2)$	$O(N^2)$
(3)	Eigendecompose $\mathbf{K}\mathbf{E} = \mathbf{E}\mathbf{v}$	$O(N^3)$	$O(N^2)$
(4)	Regularization parameter λ	$O(N^3)$	—
(5)	Estimate weights $\hat{\mathbf{c}} = \mathbf{f}(\lambda, \mathbf{y}, \mathbf{E}, \mathbf{v})$	$O(N^3)$	—
(6)	Fit values $\hat{\mathbf{y}} = \mathbf{K}\hat{\mathbf{c}}$	—	—
(7)	Estimate local derivatives, $\hat{\Delta}_{N \times P} = [\hat{\delta}_1 \quad \hat{\delta}_2 \dots \hat{\delta}_P]$	$O(PN^3)$	$O(N^2)$

Letting i, j index observations and $p = 1, 2, \dots, P$ index x variables. Steps 4-6 are followed by uncertainty estimates.

Effect of Gender on Two-Party Preference

Trump vs. Hillary with bigKRLS

Local Derivatives (dy/dx)

female ▼


Subset

(entire sample) ▼

Display

Map ▼


X Variable



Trust the Polls? Or a Poll of Polls?



House Effects, Back By Popular Demand



Simon Jackman

in HuffPost, 10/12

pollstR: the Huffington Post R API

```
library(dplyr); library(tidyr);  
library(pollstR); library(bigKRLS)
```

```
slug <- "2016-general-election-trump-vs-clinton"  
pollstR_data <- read.csv(pollstR::chart_data_url(slug))  
glimpse(pollstR_data)
```

```
## Observations: 310
```

```
## Variables: 13
```

```
## $ Trump           <dbl> 43, 39, 43, 45, 44, 49, 44,  
## $ Clinton         <dbl> 48, 46, 49, 51, 50, 47, 49,  
## $ Other           <dbl> 8, NA, 4, 4, NA, 4, 2, NA,  
## $ Undecided       <dbl> 1, 16, 5, 0, 6, NA, 6, 6, 1,  
## $ poll_id         <int> 25893, 25827, 25894, 25876,  
## $ pollster        <fctr> YouGov/Economist, Morning  
## $ start_date      <fctr> 2016-10-01, 2016-09-30, 20  
## $ end_date        <fctr> 2016-10-03, 2016-10-02, 20  
## $ sample subpopulation <fctr> Registered Voters, Likely
```

Testing for House Effects with bigKRLS


```
y <- log(pollstR_data$Clinton/pollstR_data$Trump)

t <- grep("start", colnames(X))
time_only <- bigKRLS(y , X[,t])


P <- ncol(X)
type <- grep("asked_third_party", colnames(X))
with_features <- bigKRLS(y, X[,c(type, t:P)])
```

Not So Random Errors...

Without Poll Features



With Poll Features



Election Forecasting and bigKRLS

- ▶ We model Presidential voting as a three-step process:
 - ▶ Turnout?
 - ▶ If turnout, third party or two-party?
 - ▶ If two-party, Democrat or Republican?
- ▶ Assume that third-party voters need to be modeled differently than two-party voters (fits with Perot, Gary Johnson, etc. . .)
- ▶ For our forecast, we need to:
 - ▶ Generate probability for each step
 - ▶ Model each individual-level probability
 - ▶ Split sample and predict probability of each outcome for each individuals
 - ▶ Simulate Many Elections

Variable Selection and Preprocessing

- ▶ Generate dependent variables using Bayesian measurement model
 - ▶ MCMCpack's method for mixed factor analysis
 - ▶ ANES (American National Elections Study) Jan 2016 Data (Individual Level)
 - ▶ State-Level Data on Recent Elections
- ▶ Then, model probabilities as Function of...
 - ▶ Individual Level Data: Political Preferences, Demographics...
 - ▶ State Level Data: State Demographics (Kaggle), Geolocation
 - ▶ Survey Weights

prediction with bigKRLS

```
set.seed(1234)
train <- sample(N, 800)
test <- !(1:N %in% train)


turnout_out <- bigKRLS(p_turnout[train, ],
                      Xturnout[train, ])

predict_turnout <- predict(turnout_out,
                          Xturnout[test, ])

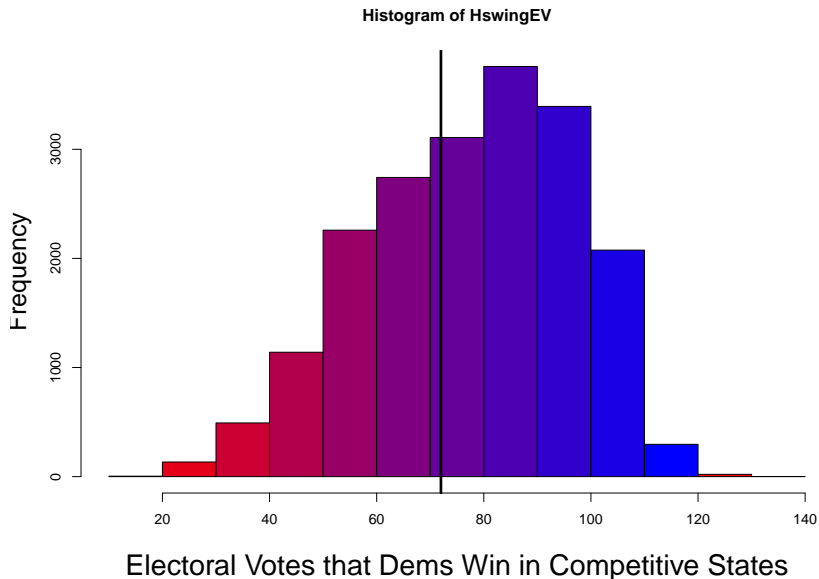
cor(predict_turnout$fit, p_turnout[test,])^2
```

Out-of-Sample Pseudo $R^2 > 0.945$ for all models

Swing State Simulations



Forecast: Hillary Wins in 62.9% of Sims



Take Aways

- ▶ *bigKRLS* is flexible and interpretable but very computationally intensive
- ▶ Not a panacea but able to replicate leading forecasts with small amounts of limited data

Next Steps for the Algorithm...

- ▶ Parallel Processing via RcppParallel or snow (in development)
- ▶ Develop practical benchmarks for assessing asymptotics
- ▶ “big” Eigentruncation (*partial_eigen* in [irlba](#) or the power method...)]
- ▶ more “Shiny” features

Thank You!!

bigKRLS on GitHub:

<https://github.com/rdr1990/bigKRLS.git>

Keep in Touch!

Robert: @rbshaffer (Me): @petemohanty